

Secrets revealed in this session:

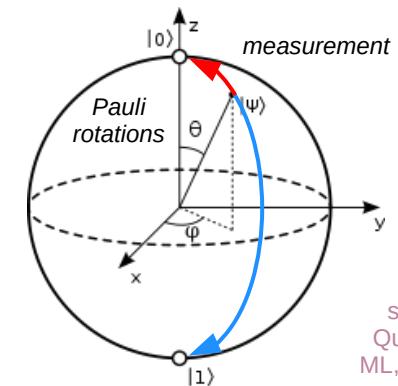
**To explore the practical aspects
of building quantum machine
learning models and their
optimisation**



QML and its aims
Parameterised circuits
Variational quantum algorithms
Data encoding / angle encoding
The good, the bad and the ugly
State measurement
Ansatz design and training
Working with DL and QML models
Model geometry and gradients
Parameters optimisation
The curse of dimensionality
QML readings
Qiskit demo and tasks (TS curve fitting)
Summary and Q&A

Quantum Algorithms and Data Encoding for QML with Qiskit

Jacob L. Cybulski
Enquanted, Melbourne, Australia



We will assume
some knowledge of
Quantum Computing
ML, Qiskit and Python

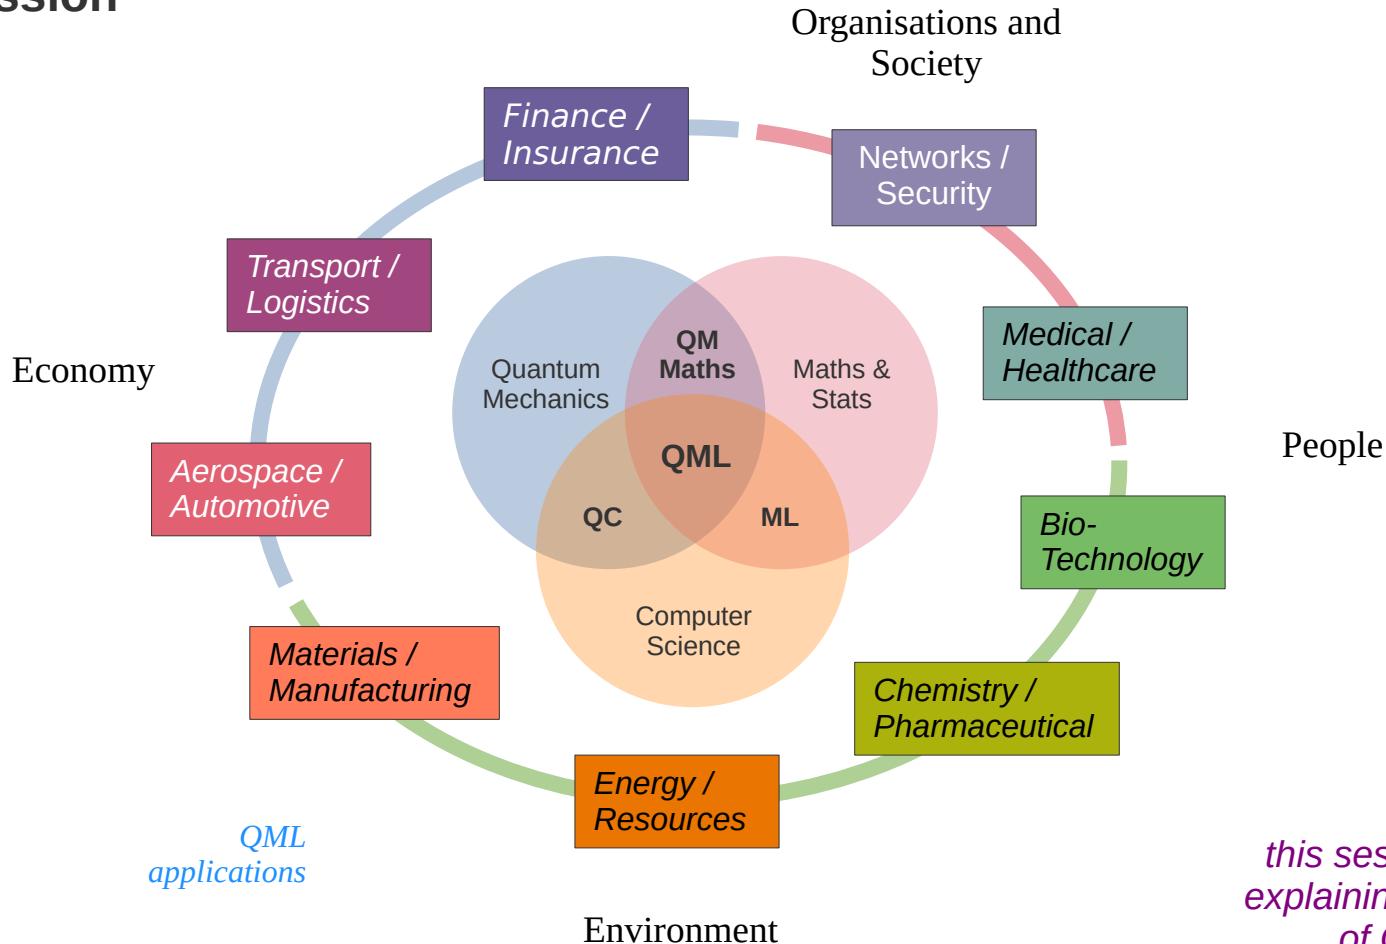
Quantum ML

aims of this session

Jacob L. Cybulski, Quantum Business Series (Deakin, RMIT, ACS, Warsaw School of Economics)
Jacob L. Cybulski, Quantum Computing Intro Series (SheQuantum, Assoc of Polish Profs in Australia)
2021-2025

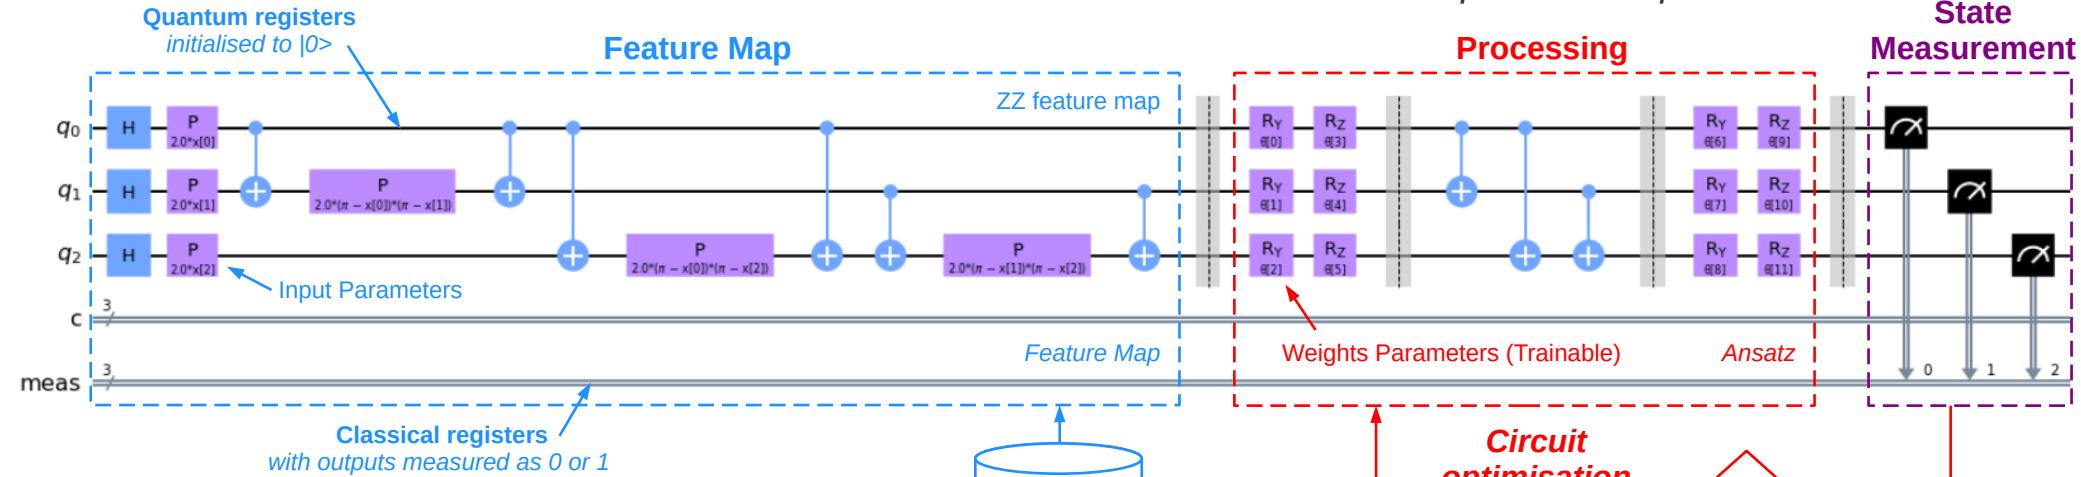


Jacob Cybulski, Founder
Enquantum, Australia



*this session aims at
explaining the nature
of QML models*

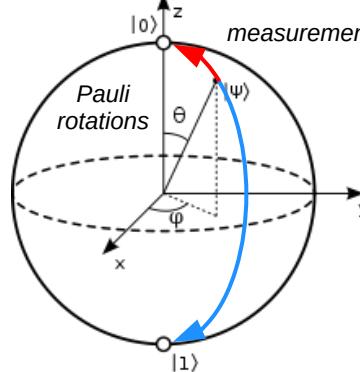
Variational Quantum Circuits and Variational Quantum Algorithms



We can create a “variational” model = a circuit template with parameterised gates, e.g. $P(a)$, $Ry(a)$ or $Rz(a)$, each allowing rotation of a qubit state in x, y or z axis (as per Bloch sphere).

Typically (but now always), the circuit consists of three blocks:

- a feature map (input)
- an ansatz (processing)
- measurements (output)



Ansatz parameters are trainable.
Each parameter defines a dimension in the model's parameter space.

Classical input data is encoded into the feature map's parameters, setting the model's initial quantum state.

The quantum state is altered by an ansatz, of parameterised gates (operations), which are trained by an optimiser

The final quantum state of the circuit is then measured and interpreted as the model's output in the form of classical data.

Circuit optimisation

cost is minimised during circuit training

Cost Fun

measured states are interpreted to match training data

Data encoding strategies

Data encoding

There are many methods of data embedding, such as:
the *basis*, *angle*, *amplitude*, *QRAM*, ... encoding,

In this workshop we will rely on *angle encoding* realised as qubit state rotation by the angle defined by the data.

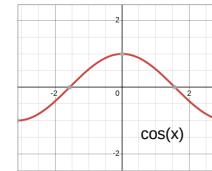
The rotation operators are always available in a quantum platform API, e.g. *Rx*, *Ry*, *Rz*, *P* or *U* (*xyz*).

Typically, the encoding rotation is performed around x or y axis, or both (allowing two values per qubit).

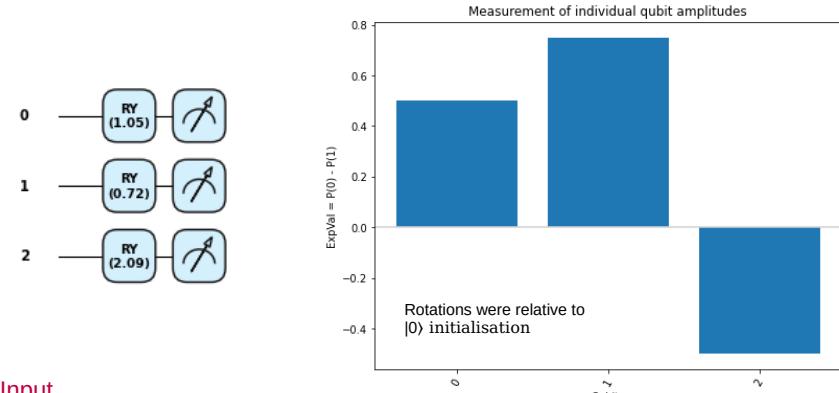
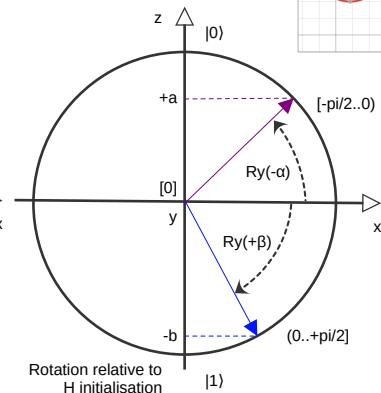
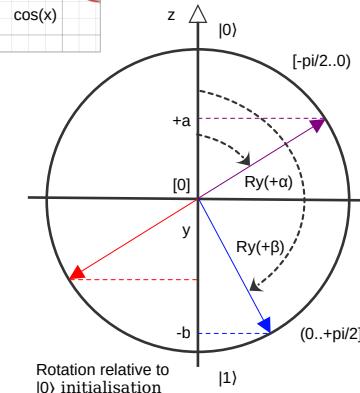
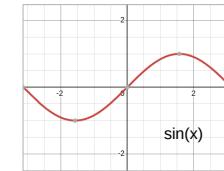
Rotations are *relative to a specific qubit state*, commonly starting at $|0\rangle$ state, or $(|0\rangle+|1\rangle)/\sqrt{2}$, which require qubits to be initialised in these states.

The encoded value could be represented either by the *angular rotation*, or the *amplitude* of the qubit projective measurement (Z).

Input data can also be repeatedly encoded and spread around the circuit, which is called *data reuploading*, and which is known to improve the model performance.



Note that training will place qubit states in areas $x < 0$ and arbitrarily around the z axis. Measurements of such states cannot distinguish them from "pure" $x > 0$ and $z = 0$.



Input

Values entered:
Ry angles used:

[np.arccos(0.5), np.arccos(0.75), np.pi-np.arccos(0.5)]
[1.047, 0.723, 2.094]

Measurements

Probabilities:
Amplitudes:

[[0.25, 0.75], [0.562, 0.438], [0.25, 0.75]]
[0.5, 0.75, -0.5]

Angle encoding

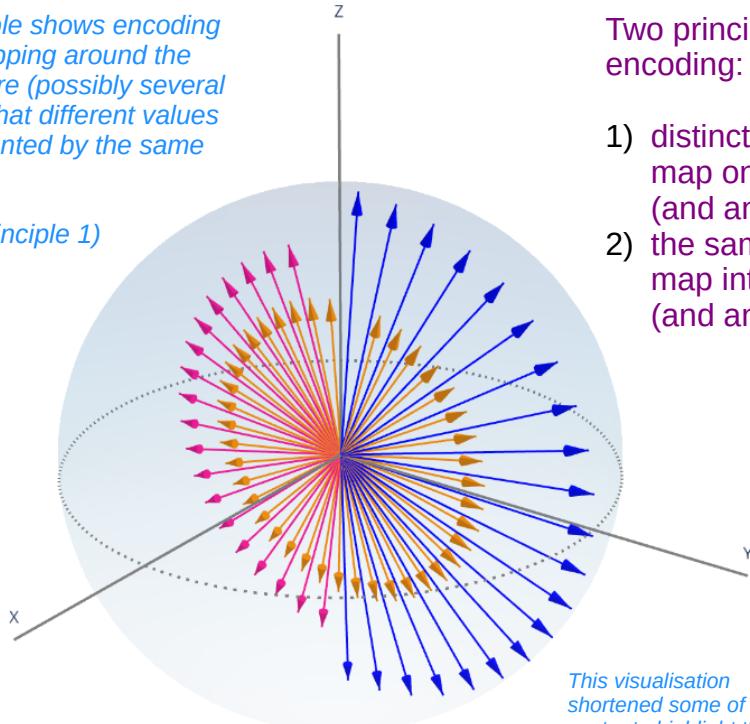
The Good, the Bad and the Ugly

These two examples explain what happens at the very beginning of model training. Later during parameters optimisation, which relies on Rx, Ry and Rz state rotations and CNOT entanglements, the model qubits potentially assume all possible states on their Bloch sphere. It is important, however, that the measured outcomes are not in conflict.

Angle Range: -6 to 6
 $0..\pi = \text{"blue"}$ (long) | $>\pi = \text{"deeppink"}$ (medium) | $-\pi..0 = \text{"darkorange"}$ (short)

This example shows encoding values wrapping around the Bloch sphere (possibly several times), so that different values are represented by the same amplitude.

(violates principle 1)



Two principles of quantum data encoding:

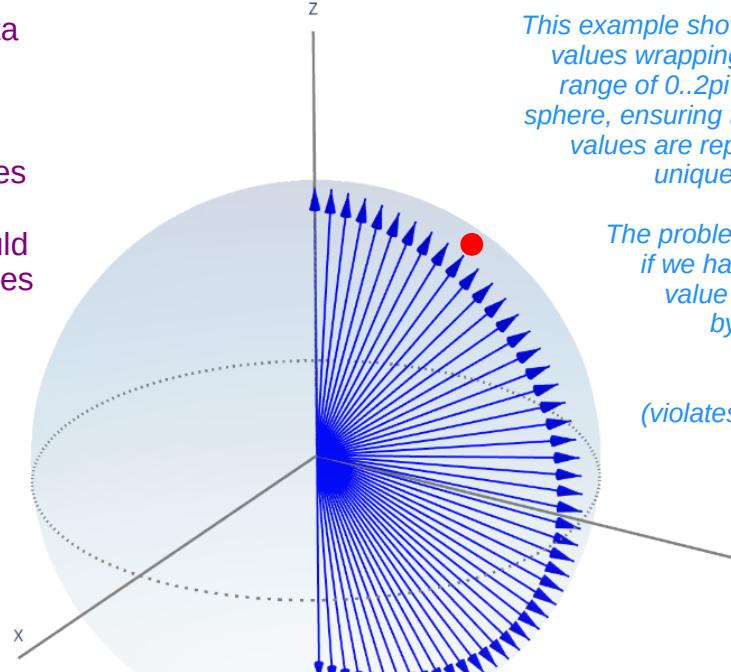
- 1) distinct data values should map onto distinct amplitudes (and angular codes)
- 2) the same data values should map into identical amplitudes (and angular codes)

This visualisation shortened some of the vector to highlight their difference, however, they are all of length 1.

Angle Range: 0 to 3.141592653589793
 $0..\pi = \text{"blue"}$ (long) | $>\pi = \text{"deeppink"}$ (medium) | $-\pi..0 = \text{"darkorange"}$ (short)

This example shows encoding values wrapping around the range of 0.2π of the Bloch sphere, ensuring that different values are represented by unique amplitudes.

The problem may arise if we have the same value represented by two distinct amplitudes
(violates principle 2)



The red dot represents the same value which on two occasions maps onto two different angles and thus two amplitudes.

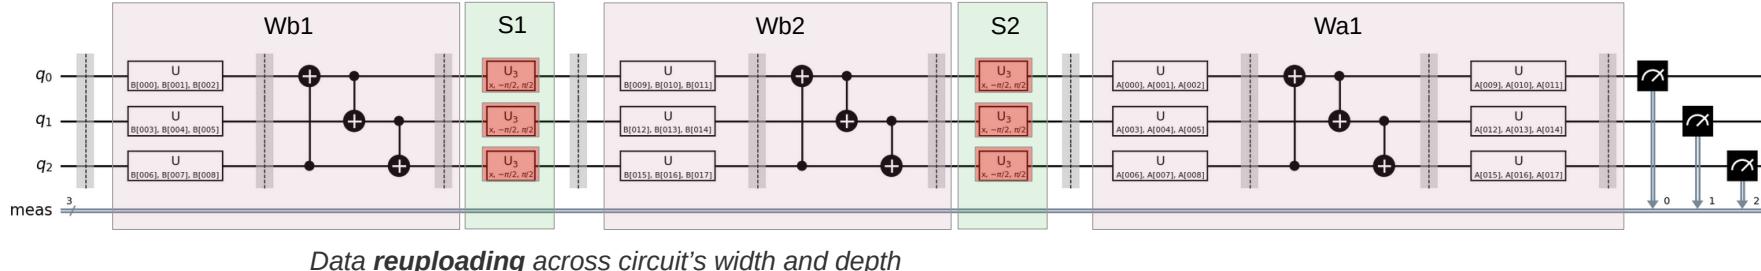
Ansatz design and training

A sample curve fitting model ...

Beware that
adding qubits adds
parameters and entanglements!

The number of states represented by the
circuit grows exponentially with the
number of qubits!

*Encoding of classical data in a quantum circuit is
not what our ML experience tells us about inputs !*



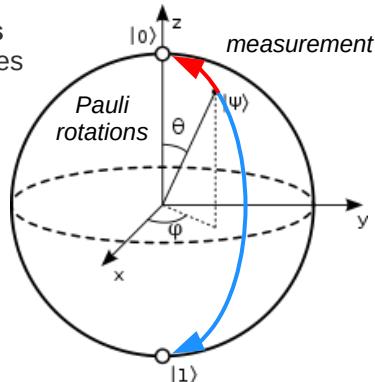
feature maps vary in:
structure and function (!!?)

ansatze vary in:

- width (qubits #)
- depth (layers #)
- dimensions (param #)
- structure (e.g. funnelling)
- entangling (circular, linear, sca)

ansatz layers consist of:
rotation blocks and entangling blocks
of $R(x, y, z)$ and CNOT gates
(rotation) (entanglement)

rotation gates
alter qubit states
around x, y, z
axes



To execute a circuit we just apply it to input data
and the optimum parameters

different cost functions:
R2, MAE, MSE, Huber, Poisson, cross-entropy,
hinge-embedding, Kullback-Leibnner divergence

different optimisers:
gradient based (Adam, NAdam and SPSA)
linear approximation methods (COBYLA)
non-linear approximation methods (BFGS)
quantum natural gradient optimiser (QNG)

circuit execution on:
simulators (CPUs), accelerators (GPUs) and
real quantum machines (QPUs)

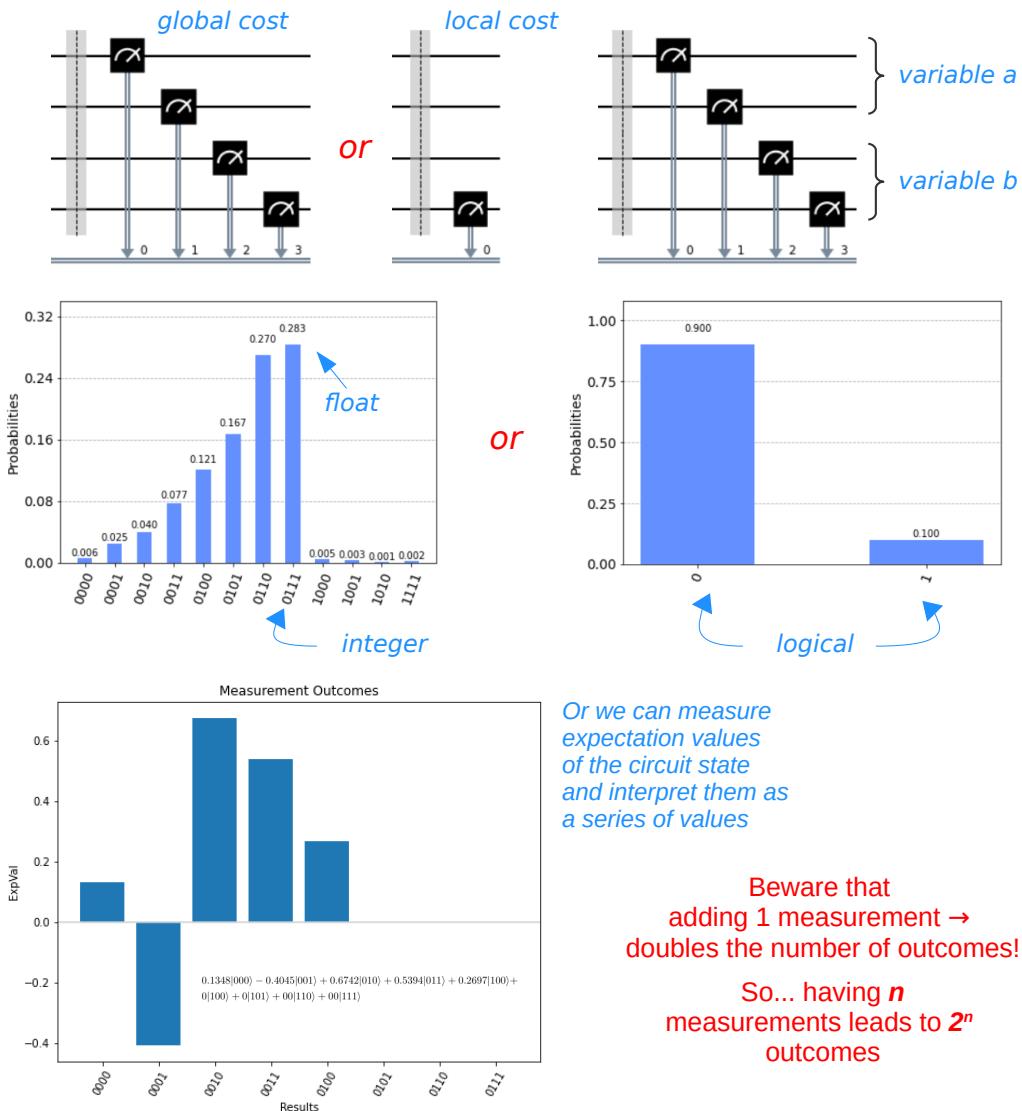
Commonly used measurements and interpretation

There are many ways of obtaining the outcome of a circuit execution, e.g. we can measure:

- all qubits (global cost / measurement)
- a few selected qubits (local cost / measurement)
- groups of qubits (each as a variable value)
- as counts of outcomes (repeated measurements)
- as probabilities of outcomes (e.g. $P(|0111\rangle)$)
- as Pauli expectation values (i.e. of eigenvalues)
- as expectation of interpreted values (e.g. 0 to 15)
- as variance, etc.

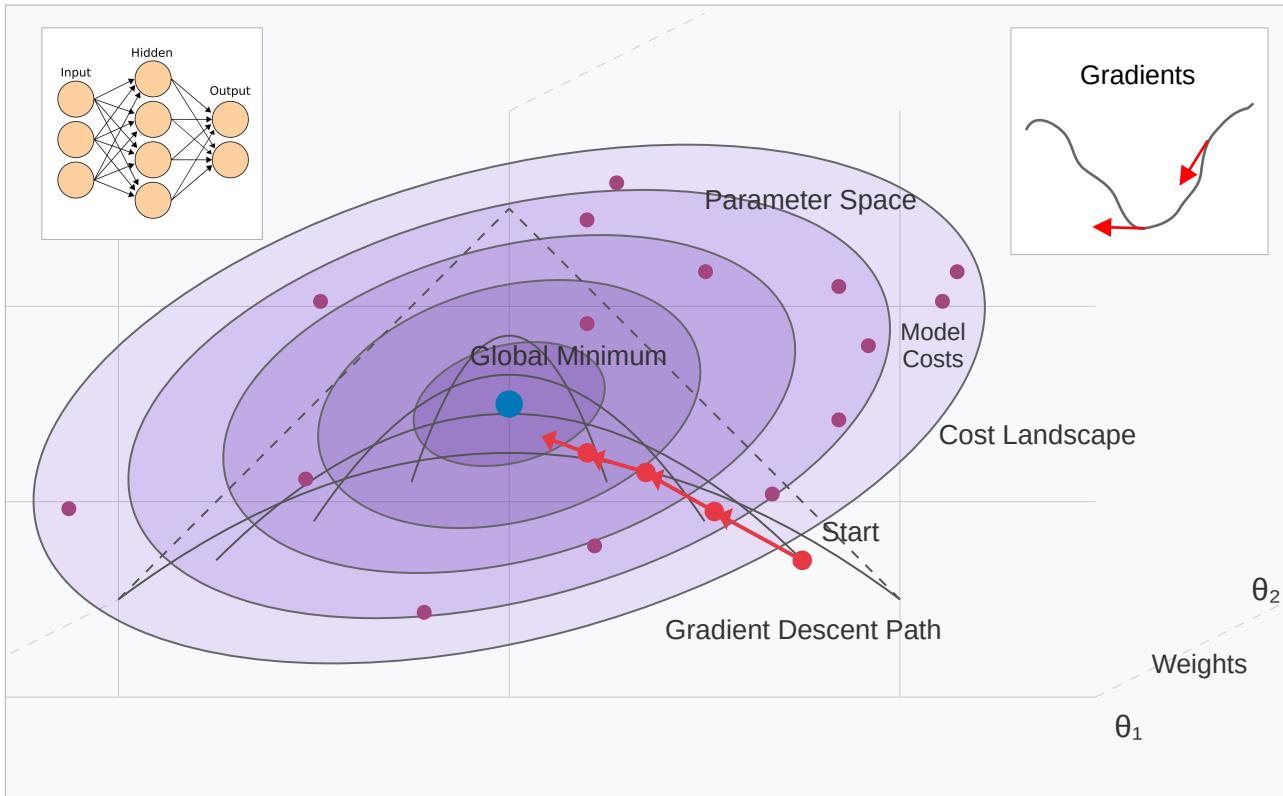
Repeated circuit measurement can be interpreted as outcomes of different types, e.g.

- as a probability distribution (as is)
- as a series of values (via expvals)
- as a binary outcome:
single qubit measurement or parity of kets
- as an integer:
most probable ket in multi-qubit measurement
- as a continuous variable:
probability of the selected ket (e.g. $|0^n\rangle$)



Working with deep learning models

Classical model optimisation



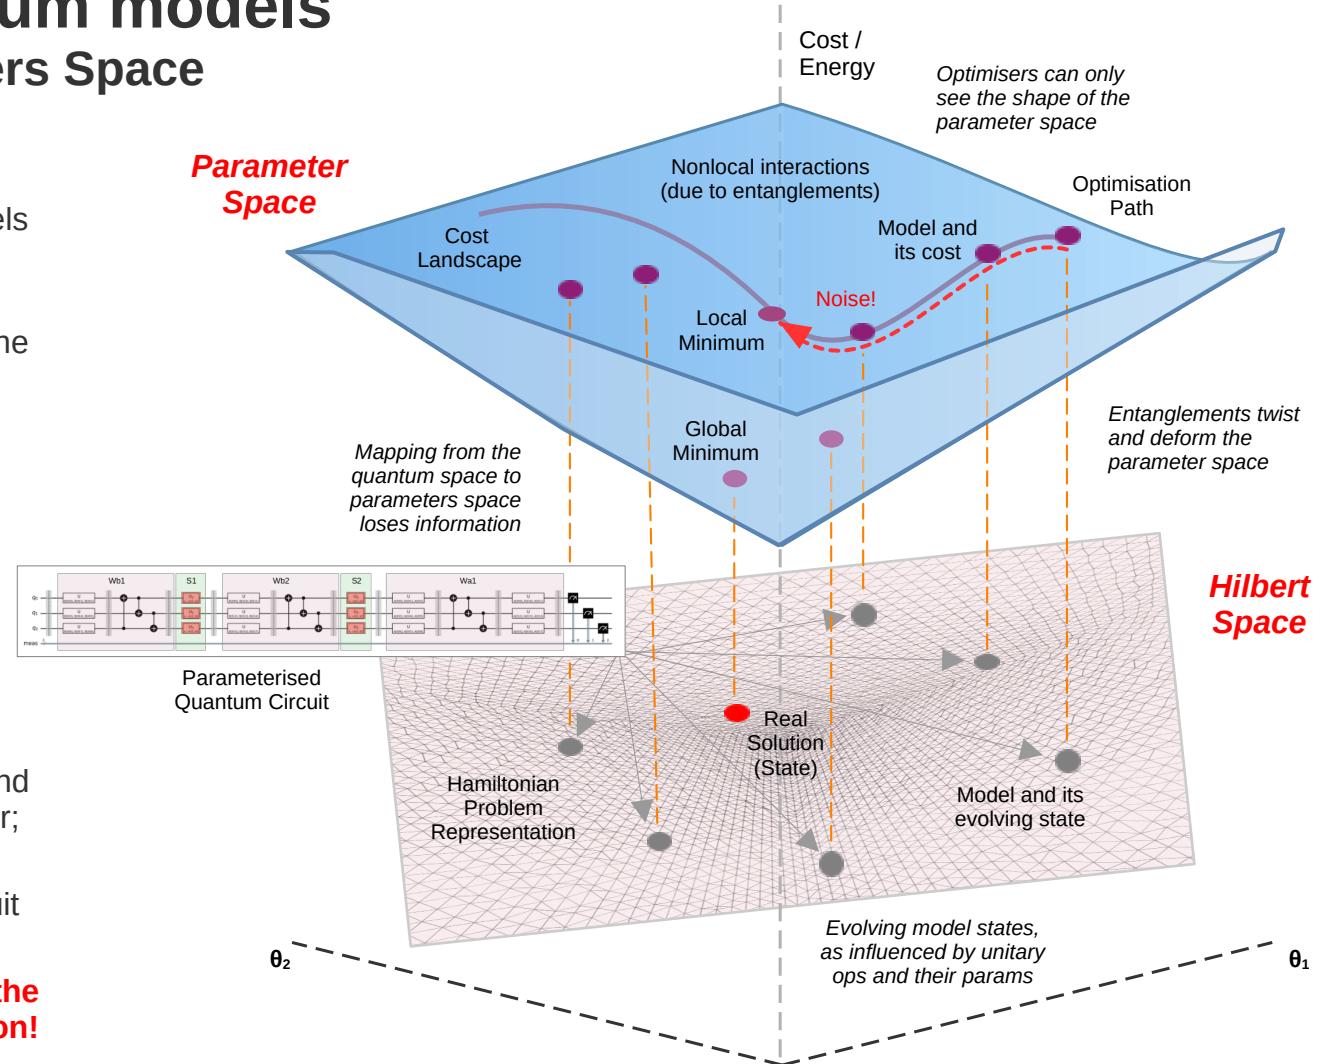
Gradients are local, i.e. their changes influence only their immediate neighbourhood

- Deep learning (DL) models evolves layers of consecutive *representations of data on input*
- *Weights* and *biases* parameterise DL models
- Model *parameters are independent*, with *local interaction* only, limited to neighbourhood
- *Cost function* evaluates the model quality, where the lower the cost, the better the model
- Differently parameterised models form a multi-dimensional geometry = *the parameter space*
- The costs of the all DL models form a surface of this geometry = *the cost landscape*
- *Optimisation algorithms* navigate the the cost landscape in search for the minimum cost
- This process may rely on the shape of the cost landscape and its *gradients*
- *Gradient descent* algorithms navigates gradients towards the *minimum cost* model
- *Backpropagation* implements gradient descent very efficiently by relying on the knowledge of local interactions between model parameters
- In the *forward step*, a gradient-based optimiser calculates the activation of DL model nodes, producing the model's output
- In the *backward step* it backpropagates its calculation error, and recalculates DL model weights and biases to correct the error

Working with quantum models

Hilbert Space vs Parameters Space

- **Hilbert state space** (dim = the number of qubits) is the quantum realm where the models and their states evolve in response to unitary operations as defined by the circuit gates;
- **Data encoding** brings in classical data into the Hilbert space as unique and correlated quantum states during the model execution;
- **Layers of circuit gates** determine the evolution of the quantum model's initial state into its final state;
- **Trainable parameter space** is a classical multi-dimensional space of circuit gate parameters, which the optimiser navigates;
- **Entanglements** (defined by CNOTs) create and correlate non-separable qubit states, which alter the parameter space geometry, and also the cost landscape used by the optimiser;
- **Measurement** of individual qubits collapses their states, consequently projecting the circuit state onto classical outcomes.
- **The mapping from the quantum space to the classical parameter loses some information!**

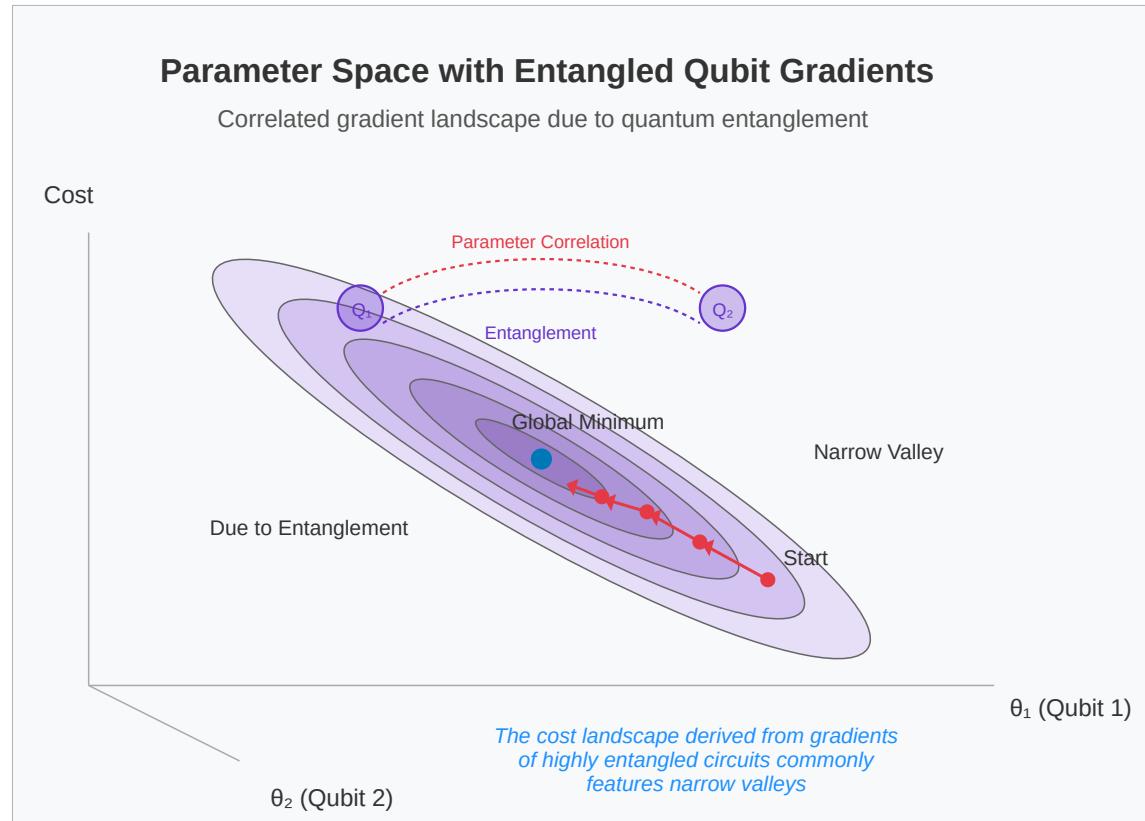


- All models generated from a PQC can be defined by the possible valid combination of its parameters
- A set of all valid models form some *geometry* in quantum state space and classic parameter space

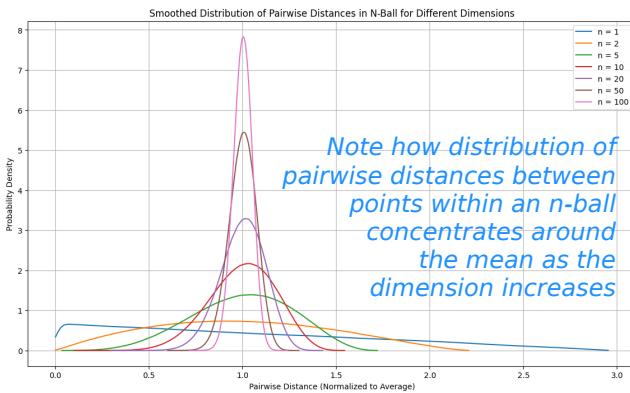
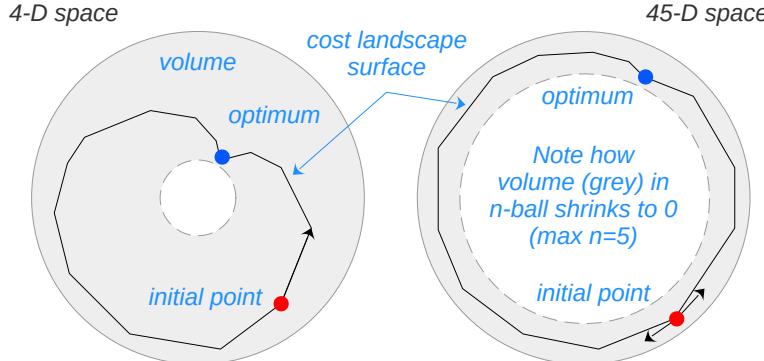
- This geometry is commonly combined with a set of model observables, which create a *navigable surface*, e.g. a quantum *energy landscape* or classical *cost landscape*
- *Classical optimisers* can then be used to navigate the classical parameter space in search for the model configuration associated with the *minimum cost*
- Such optimisation of quantum models needs unique approaches due to the presence of *non-local gradients*
- *Entangled qubits* result in *correlated parameters*, and indirectly gradients, so the changes to one are reflected in the changes to distant others
- *Backpropagation* cannot thus be used effectively in training quantum circuits, as the model states, and thus local gradients, are not directly accessible and their measurements collapse the state
- *Gradient descend* can still be used with *global gradients*, derived from the geometry of the cost landscape rather than from the relationships between parameters
- Other optimisation techniques can be used instead of gradient-based methods, e.g. *stochastic optimisation* or *particle swarm optimisation*

Quantum model optimisation

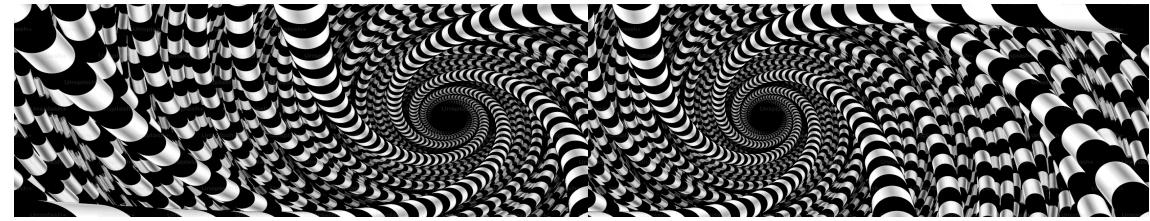
model geometry vs. classical optimisers



The curse of dimensionality



Cybulski, J.L., Nguyen, T., 2023. "Impact of barren plateaus countermeasures on the quantum neural network capacity to learn", Quantum Inf Processing 22, 442.



Barren Plateaus (too many parameters)

- Pairwise distances between uniformly distributed points in high-dimensional space become (almost) identical, and the surface of such a space is almost flat (n -ball value is near its surface).
- In a quantum model with a high-D parameter space, the cost landscape is nearly flat, the situation called **barren plateau (BP)**.
- In high-D parameter space, models sampled by the optimiser are very sparse in both Hilbert space and parameter space.
- When BPs emerge, the optimiser struggles finding the optimum.
- Selecting the optimisation initial point far from the optimum (e.g. random) makes it even more difficult !

There are some well-known BP countermeasures

- use fewer qubits / layers / parameters
- use local cost functions (do not measure all qubits)
- use non-Euclidean metrics (e.g. Fisher Information Metric)
- beware of random params initialisation (and keep them small)
- use BP-resistant model design (e.g. layer-by-layer dev)
- use BP-resistant models (e.g. QCNNs)

Quantum model dev. training, validation and testing

While developing a quantum model we need to test its capacity to learn. This can be done by applying the model to data used in its training, with the objective to test its ability to recall what it learnt.



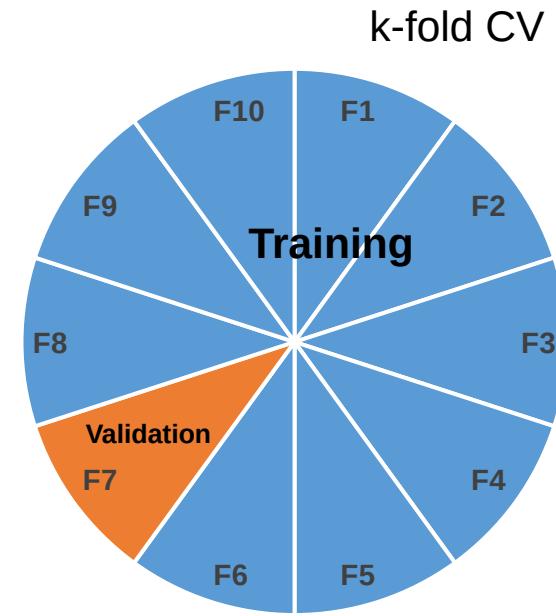
To assess the model's performance on new data we need to apply it to data not used in training.

Commonly adopted approach is known as *holdout testing*, where we randomly split data into three partitions - one to be used for model *training*, one for its *validation* while improving the model, and one for *honest testing* on data never used before.

Holdout testing validates and test the model once only, assuming all data partitions as representative of the population, which may not be true.

Cross-validation (CV) is a more appropriate testing method, which repeatedly trains and validates a model using different data samples (folds), then averaging the performance obtained from all runs.

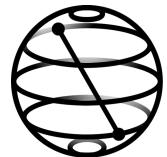
As training a quantum model may be extremely slow, cross-validation may not be feasible.



As quantum models are sensitive to their initialisation, once the model is developed, we retest it with differently initialised weights, in the process also resampling data partitions.

Note: as we develop the model we freeze data partitions by setting a *random seed*, to ensure performance changes are due to our actions and not differences in data partitions.

Qiskit QML Workshop



Qiskit ...

- Accessible from *Python*, *Rust*, *C++* and more...
- Has a standard set of *quantum state operations*
- Supports creation of flexible QML *algorithms*
- Executes on *simulators* and *quantum hardware*
- Supports hardware *accelerators* (e.g. *GPUs*)
- Provides tools for *error mitigation*
- Utilises variety of *quantum gradients models*
- Supports *hybrid quantum-classical models*
- Provides many QML models, e.g. *QNNs*, *QCNN*, *QAE*, *QSVM* and *Bayesian models*
- Can be extended with *PyTorch* and *TensorFlow*
- Among quantum SDKs, it is *the best performer*
- It is largely *hardware agnostic via vendor backends*
- Supports *IBM quantum backend and runtime*
- It is *complex* and its *core design changes too often!*

Qiskit QML tasks (time series curve fitting):

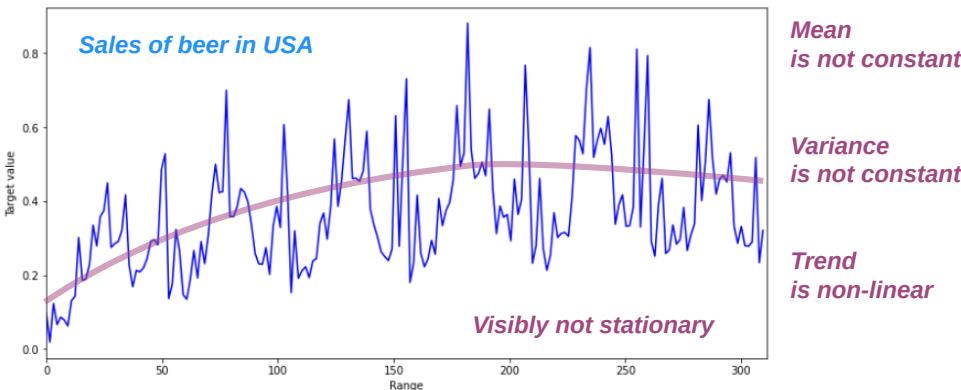
- Add ML 0.8.3 package to Qiskit 1.4.4 (Python 3.11)
- Create a few simple models to fit samples of data
- Learn to initialise model weights
- Learn the interaction b|n data encoding and ansatz
- Understand how observables / measurements work
- Explore the impact of ansatz structure on performance
- **Challenge:** Apply your skills to complex data
- **Reflection:** Refine your QML development process

Key takeaways:

- Plan model development, tests and experiments
- More params and entanglements improve *expressivity*
- Data reuploading makes a huge difference!
- More width / depth / params reduce *trainability* = *the curse of dimensionality*
- More entanglements reduce *trainability*
- High dimensional param space upsets even non-gradient optimisers due to *model sparsity*
- Bad data encoding spoils the bunch!
- Carefully consider your quantum model initialisation
- Surprise - a single qubit model still works! (and well)
- More training often does not eliminate problems!

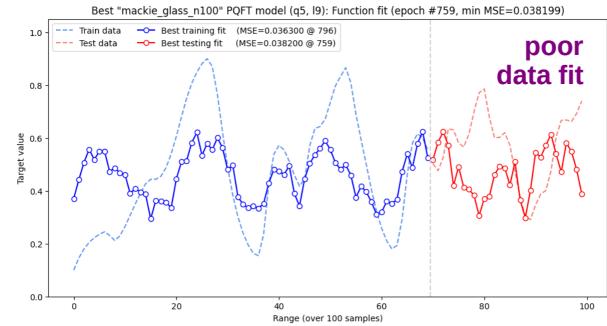
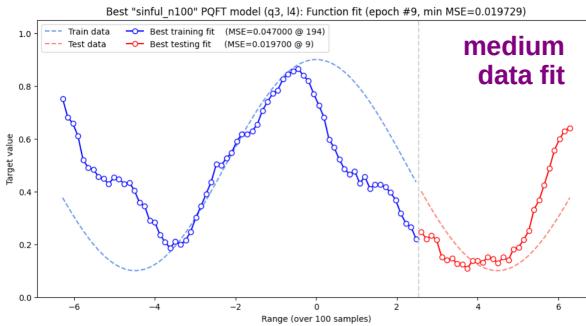
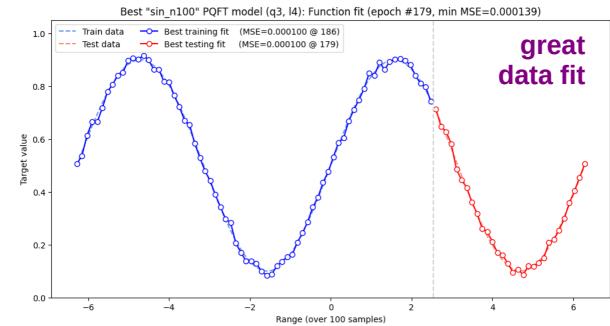
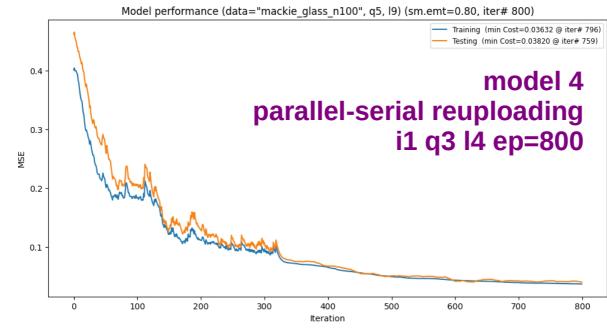
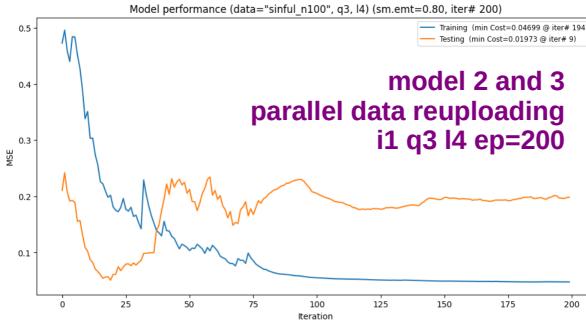
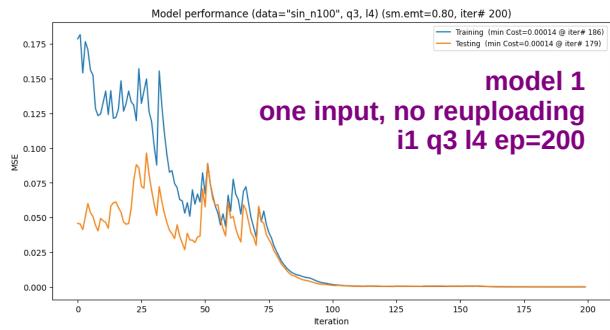
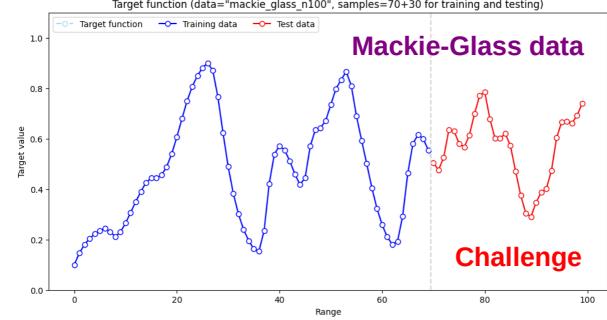
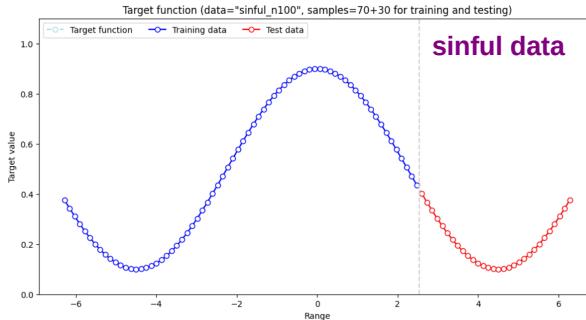
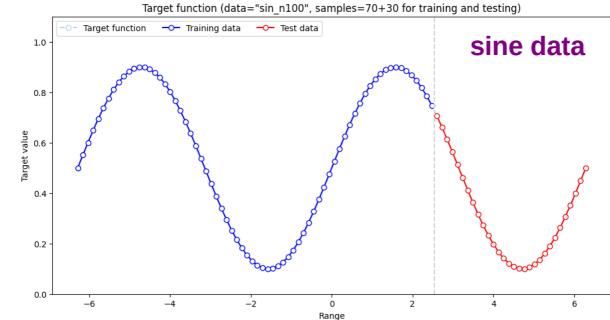
Key concepts in time series analysis

- Time series (TS) analysis aims to *identify patterns* in historical time data and to *create forecasts* of what data is likely to be collected in the future
- *Many TS applications*, including heart monitoring, weather forecasts, machine condition monitoring, etc.
- *Many excellent tools* and *efficient methods* of TS analysis, yet improvements are still desirable
- Time series must have an *unique temporal index* e.g. a time-stamp sequencing the series
- Time series must also have some *time-dependent attributes* suitable for modelling



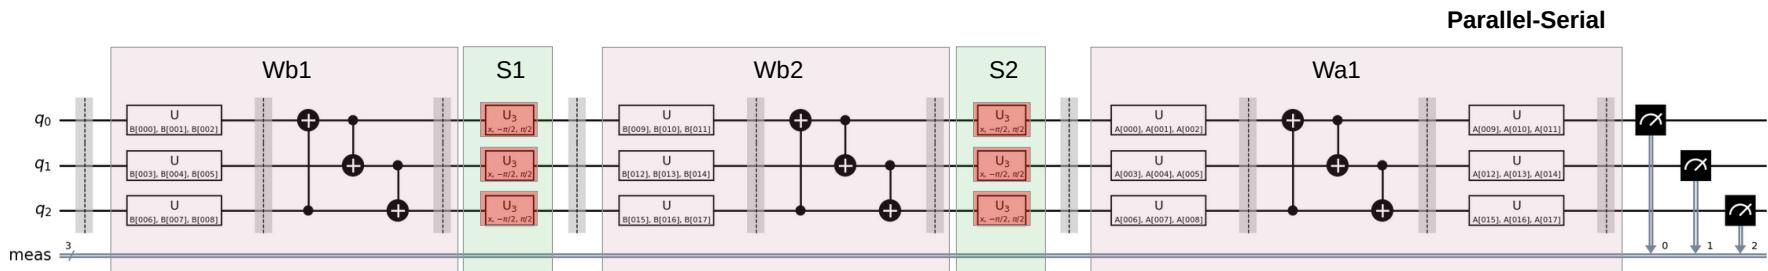
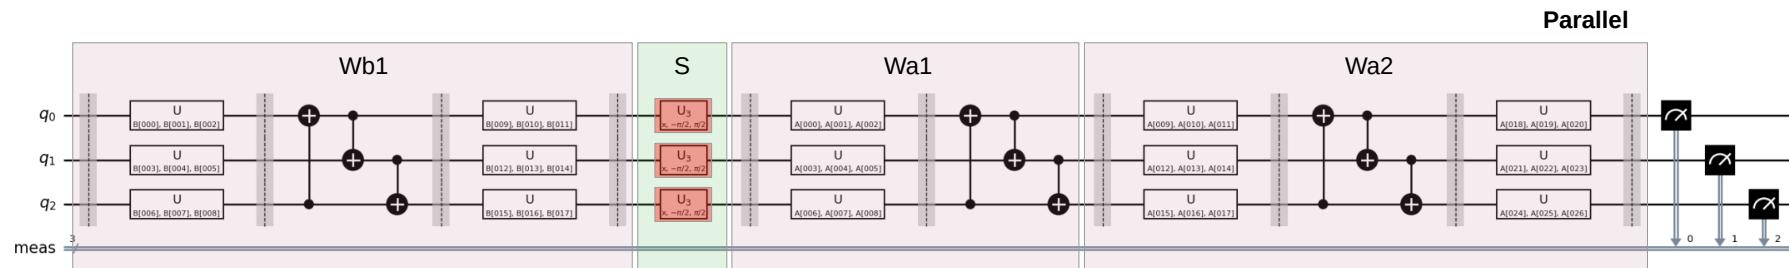
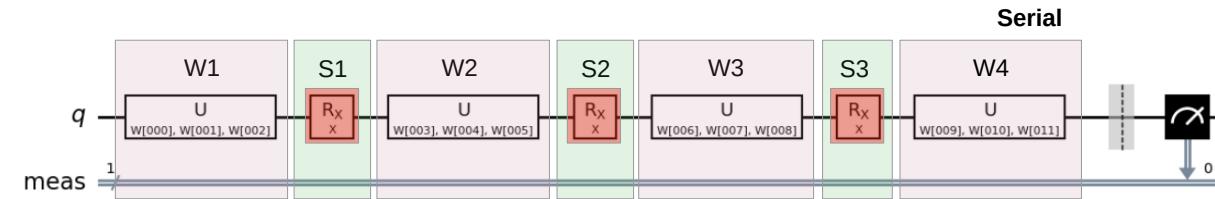
- Time series can be *univariate* or *multivariate*
- Time series often show *seasonality* in data, i.e. some patterns repeating over time
- **TS values are dependent on the preceding values!**
- **Distinction between consecutive TS values is small!**
- There are several different types of TS models, e.g.
 - The first group are *curve-fitting models*, which are trained to fit a function to a sample of data points, to predict data values at specific points in time
 - The second group are *forecasting models*, which are trained to predict future data points from their preceding temporal context (a fixed-size window sliding over TS)
- Majority of statistical forecasting methods require *strict data preparation*, i.e. dimensionality reduction, TS aggregation, imputation of missing values, removal of noise and outliers, adherence to normality and homoskedasticity (e.g. they need to be stationary)
- **QML methods do not have such strict requirements, and are promising for effective time series analysis and forecasting!**

Data and some results



Tasks: improve curve-fitting of these models

Curve-fitting models



Qiskit models

Custom circuit creation

```
##### Model 1 - Single input
def qnn_model_1(qubits_no, layers_no, add_meas=False):
    qr = QuantumRegister(qubits_no, 'q')
    ansatz = QuantumCircuit(qr, name="ansatz")
    param_x = Parameter('X')
    ansatz.rx(param_x, 0)

    for l in range(layers_no):
        ansatz.barrier()
        for q in range(qubits_no):
            ansatz.rx(Parameter(f'P[{l}:02d]_A{q:02d}'), q)
            ansatz.ry(Parameter(f'P[{l}:02d]_B{q:02d}'), q)
            ansatz.rz(Parameter(f'P[{l}:02d]_C{q:02d}'), q)
        for q in range(qubits_no-1):
            ansatz.cx(q, q+1)
        if qubits_no > 1:
            ansatz.cx(qubits_no-1, 0)
    if add_meas:
        ansatz.measure_all()
    return ansatz
```

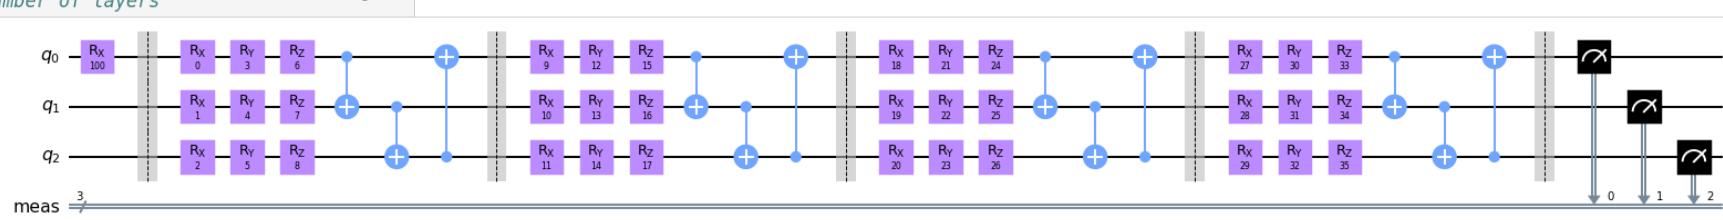
define a quantum model 1

```
##### Model settings for models with 1 input
```

```
inputs_no = 1      # Number of inputs
qubits_no = 3      # Number of qubits
layers_no = 4      # Number of layers
epochs = 200       # Tr
shots = 1000       # Ig
print_fract = 0.1  # Fr
weights_scaler = 1.00 # Sc
```

model circuit

2
and its settings



Here we present a simple curve-fitting quantum model developed in Qiskit.

The model's architecture consists of a feature map of a single Rx gate and several layers (4), each layer providing trainable parameters (Rx, Ry, Rz) spanning all qubits (3 x 3), as well as, an entangling block of CNOT gates arranged in a circular fashion (3).

The following code illustrates model creation and its preparation for training:

- [1] Function creating parameterised model circuit (inputs_no, qubits_no, layers_no).
- [2] Meta-parameters defining the model and its training (epochs, shots, weight_scaler).
- [3] Selection of the model function, loss function and optimiser.
- [4] Creation of the circuit and a vector of initial weight values.

Select which model to use

```
mfun = qnn_model_1
```

Select the loss function (here, MSE cost)

loss_fun = L2Loss()

Select the optimiser (here, non-gradient based)

optimizer = COBYLA(maxiter=epochs)

3
select:

- model
- loss function
- optimiser

algorithm_globals.random_seed = seed
np.random.seed(seed)

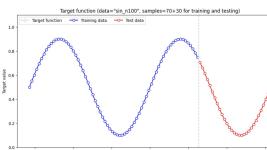
Create and initialise a quantum model
model = mfun(qubits_no, layers_no)
init_weights = weights_scaler * 2 * np.pi * \\\nalgorithm_globals.random(model.num_parameters - inputs_no)

Define a QNN estimator

obs = SparsePauliOp.from_list([(Z ** model.num_qubits, 1.0)])

4
create model circuit
calculate initial weights

Sample data needs to be prepared, cleaned and split into training and test partitions.



Training a simple Qiskit estimator

```
model = mfun(qubits_no, layers_no)
init_weights = weights_scaler * 2 * np.pi * \
    algorithm_globals.random(model.num_parameters - inputs_no)

### Define a QNN estimator

obs = SparsePauliOp.from_list([('Z' * model.num_qubits, 1)])
estimator = StatevectorEstimator(seed=seed)
regr_qnn = EstimatorQNN(
    circuit=model,
    input_params=[model.parameters[-inputs_no]],
    weight_params=model.parameters[:-inputs_no],
    observables=obs,
    estimator=estimator,
    gradient=ParamShiftEstimatorGradient(estimator),
)
```

```
### Define and fit regressor

regr_callback = create_callback(epochs, print_fract=print_fract)

regressor = NeuralNetworkRegressor(
    neural_network=regr_qnn,
    loss=loss_fun,
    optimizer=optimizer,
    initial_point=init_weights,
    callback=regr_callback
)
    # Select the loss function (here, MSE cost)
    loss_fun = L2Loss()
    # Select the optimiser (here, non-gradient based)
    optimizer = COBYLA(maxiter=epochs)

### Train the model
print(f'\nModel training started\n')

    callback=regr_callback
)
    # Select the loss function (here, MSE cost)
    loss_fun = L2Loss()
    # Select the optimiser (here, non-gradient based)
    optimizer = COBYLA(maxiter=epochs)

### Train the model
print(f'\nModel training started\n')

    regressor.fit(X_train, y_train)
    elapsed = time.time() - start
    time_str = time.strftime("%H:%M:%S", time.gmtime(elapsed))
```

training loop 7

```
Model training started          training log
(00:00:00) - Iter#:  0 / 200, Cost: 0.178608
(00:00:01) - Iter#:  20 / 200, Cost: 0.176996
(00:00:02) - Iter#:  40 / 200, Cost: 0.028761
(00:00:03) - Iter#:  60 / 200, Cost: 0.008521
(00:00:05) - Iter#:  80 / 200, Cost: 0.004425
(00:00:06) - Iter#: 100 / 200, Cost: 0.001005
(00:00:07) - Iter#: 120 / 200, Cost: 0.000322
(00:00:09) - Iter#: 140 / 200, Cost: 0.000784
(00:00:10) - Iter#: 160 / 200, Cost: 0.000176
(00:00:11) - Iter#: 180 / 200, Cost: 0.000145

Total time 00:00:12, min Cost=0.000144
```

5 estimator

6 regressor

In Qiskit it is possible to create variety of different quantum models.

Here we will show to build a custom circuit using a Qiskit state vector estimator and a neural network regressor.

```
##### Callback function use to collect training data

objfun_vals = [] # To store objective function values
params_vals = [] # To store parameter values

def create_callback(epochs, print_fract=0.1):

    global objfun_vals, params_vals
    objfun_vals = []
    params_vals = []
    elapsed = 0
    start_time = time.time()

    def callback_func(weights, obj_func_eval):
        nonlocal epochs, print_fract, start_time, elapsed
        global objfun_vals, params_vals

        iters = len(objfun_vals)
        objfun_vals.append(obj_func_eval)
        params_vals.append(weights)
        elapsed = time.time() - start_time
        time_str = time.strftime("%H:%M:%S", time.gmtime(elapsed))
        if (print_fract == 0) or (iters % int(print_fract*epochs) == 0):
            print(f'{time_str} - Iter#: {iters:3d} / {epochs:3d}, '+
                  f'Cost: {obj_func_eval:.6f}')

    return callback_func
```

```
regressor.fit(X_train, y_train)
elapsed = time.time() - start
time_str = time.strftime("%H:%M:%S", time.gmtime(elapsed))

### Retrieve training cost and params as arrays
objfun_vals = np.array(objfun_vals)
params_vals = np.array(params_vals)
```

callback function

8

Qiskit ML provides a utility *function "fit"* which executes a *training loop* [7], folding into one: a *forward* step which applies the model with its current parameters to training data, *cost calculation*, and a *backward* step to improve the model parameters.

callback info

Quantum model performance: Scoring your quantum model

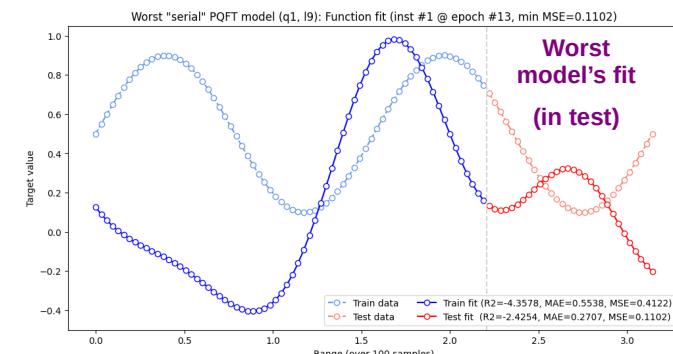
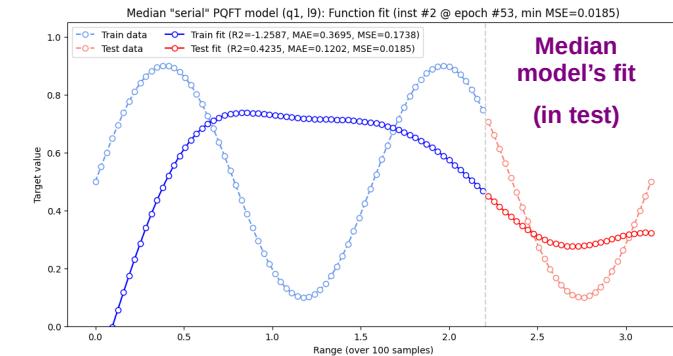
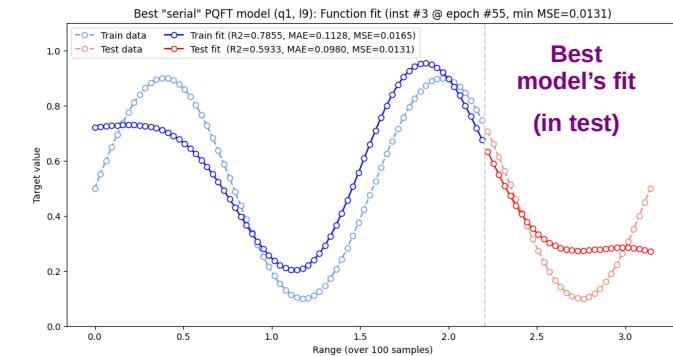
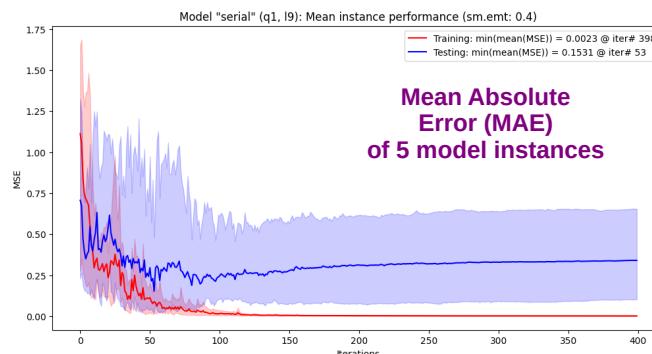
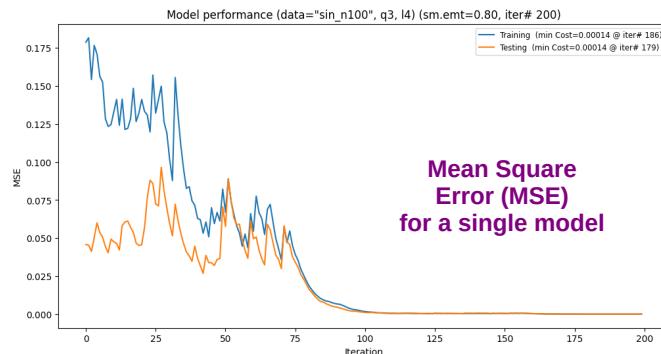
In model training we rely on the specific loss function to guide the optimiser, e.g. L2Loss (MSE cost).

In training, the costs and the model parameters for all optimisation steps are saved. Later, the parameters values are assigned to the weights of the circuit, which is then applied to both training and test data, and results scored using variety of metrics (e.g. MAE and R^2).

However, as a quantum model performance is highly sensitive to its initialisation, it is also advisable to run multiple, differently initialised, instances of the same model. Subsequently analyse distribution of their performance results (here we present only 5 instances of the same model in identical meta-parameter configurations).

When doing so, it is also possible to present the level of model's fit to data, depending on its best, median or worst instance performance.

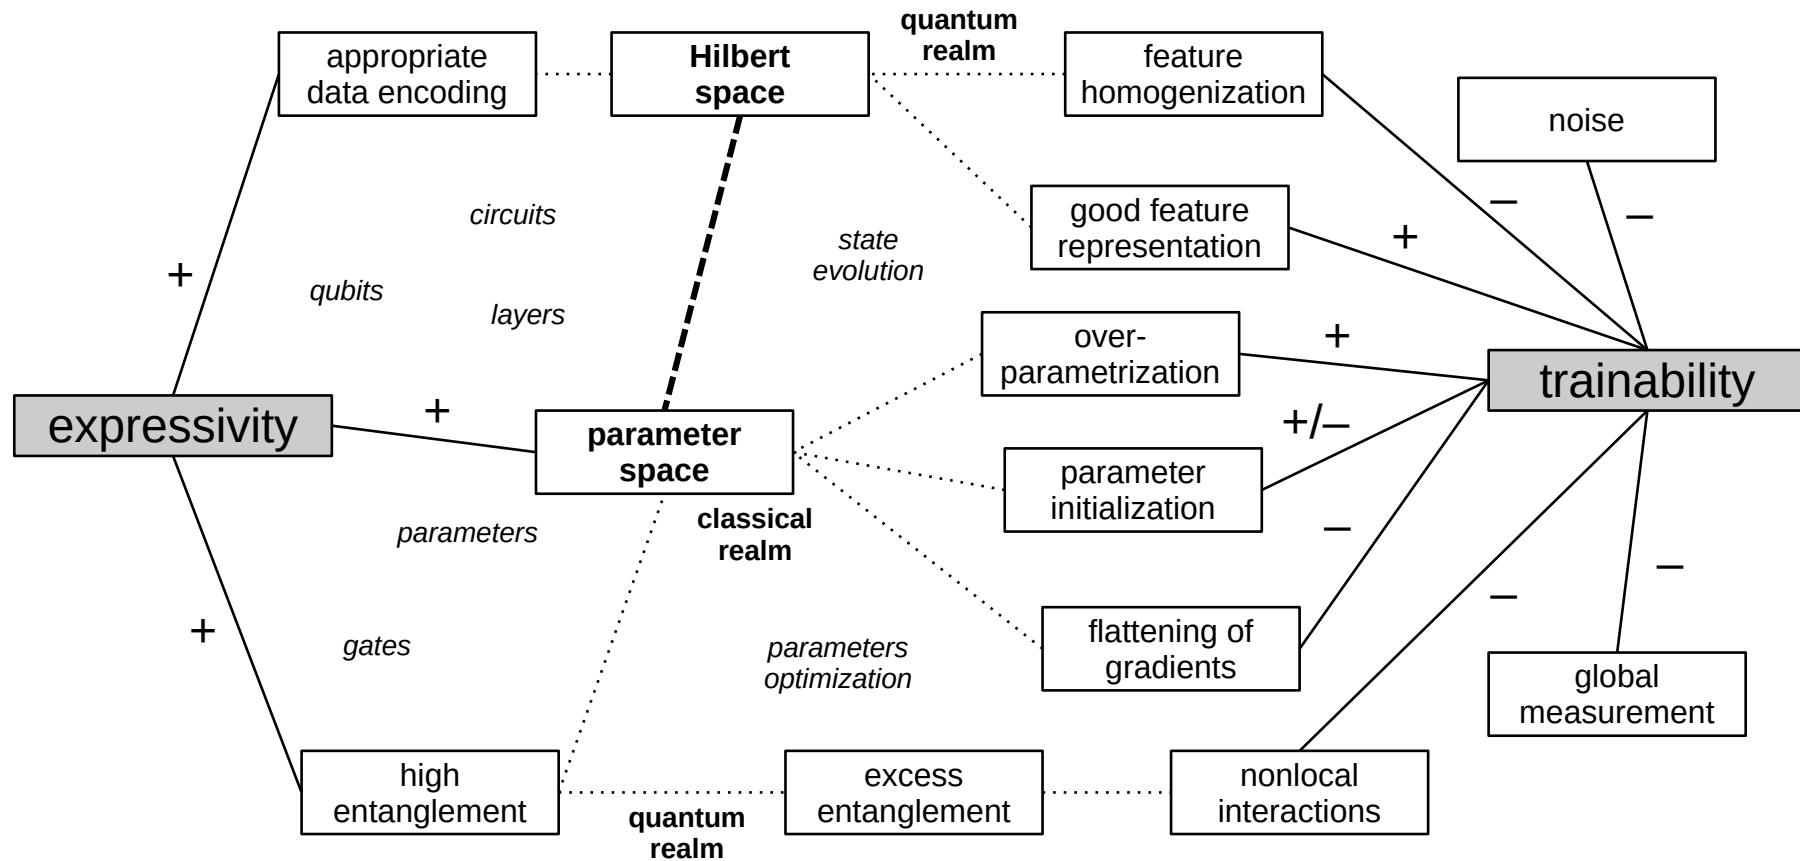
In this way, our performance assessment is not subject to our luck or the lack thereof, and most importantly we can provide an honest and unbiased reporting on the model's fitness for its practical deployment.



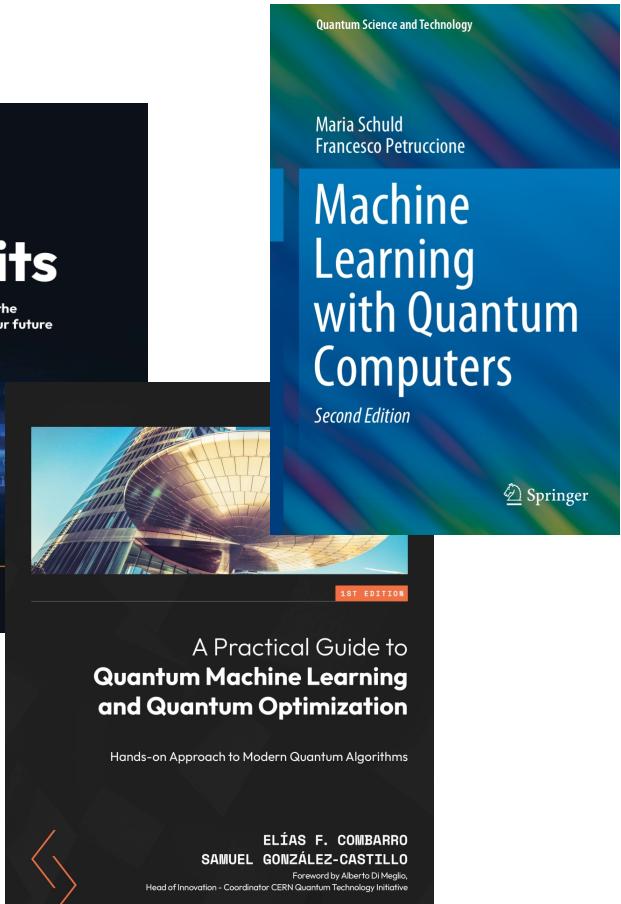
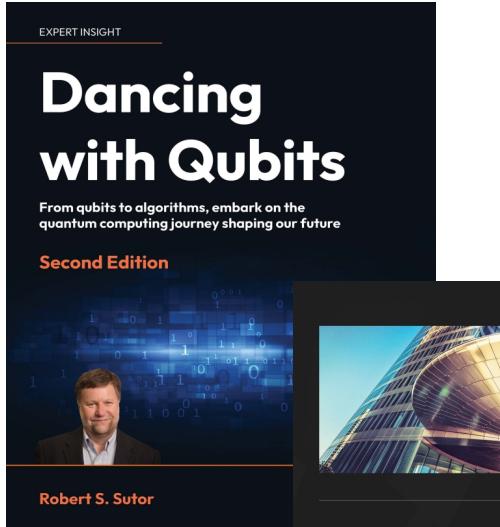
Expressivity vs. Trainability

Model expressivity: ability to effectively represent time series data in quantum space.

Model trainability: capacity to learn and generalise for predictive accuracy and efficiency in the process of model optimisation.



Recommended reading on QML with Qiskit



Quantum computing with Qiskit

Ali Javadi-Abhari,¹ Matthew Trinastic,¹ Kevin Krissian,¹ Christopher J. Wood,¹ Jake Lishman,¹ Julian Gacon,³ Simon Martis,⁴ Paul D. Nation,¹ Lev S. Bishop,¹ Andrew W. Cross,¹ Blake R. Johnson,¹ and Jay M. Gambetta¹

¹*IBM Quantum, IBM T.J. Watson Research Center, Yorktown Heights, NY, 10598*
²*IBM Quantum, IBM Research Europe, Hursley, United Kingdom*
³*IBM Quantum, IBM Research Europe, Zürich, Switzerland*
⁴*IBM Quantum, IBM France Lab, Orsay, France*

We describe Qiskit, a software development kit for quantum information science. We discuss the key design decisions that have shaped its development, and examine the software architecture and runtime system. We also highlight the quantum circuit abstraction, which is the core building block of matter physics on a quantum computer that serves to highlight some of Qiskit's capabilities, for example the representation and optimization of circuits at various abstraction levels, its scalability and reusability to new gates, and the use of quantum-classical computations via dynamic circuits. Lastly, we discuss some of the ecosystem of tools and plugins that extend Qiskit for various tasks, and the future ahead.

I. INTRODUCTION

Quantum computing is progressing at a rapid pace, and robust software tools such as Qiskit are becoming increasingly important as a means of facilitating research, education, and to run computationally interesting problems on quantum computers. For example, one of the first applications of Qiskit was the optimization of qubit mapping [51], a bit-tolerant identity [42], inspired of qiling gates [101, 104], or the circuit editor toolbox after [1]. The packt is a collection of books have been created by whom the Python as well as human [3, 4]. More have used and used Qiskit is the most [10]. A milestone in philosophy drive quantum circuit form capability function a quantum lever a very work and pressio system end.

II. DESIGN PHILOSOPHY

We begin by discussing Qiskit's scope within the broader quantum computing software stack, as illustrated in Figure 1. Starting from a computational problem, a quantum algorithm specifies how the problem must be solved on a quantum computer. This step involves translating the classical problem to the quantum domain, for example Fermion to qubit mapping [31, 62]. Circuits at this level can be quite abstract, for example only specifying a set of Pauli rotations, some unitaries, or other high-level mathematical operators. Importantly, these abstract circuits are representable in Qiskit, which contains synthesis methods to generate concrete circuits from them. Such concrete circuits are generated using a standard library of gates, represented using intermediate language constructs such as OpenQASM [2].

The transpiler translates circuits in multiple rounds of passes, in order to optimize and translate it to the target instruction set architecture (ISA). The word "transpiler" is used within Qiskit to emphasize its nature as a circuit-to-circuit rewriting tool, distinct from a full compilation down to controller binaries which is necessary for executing circuits. But the transpiler can also be thought of as an optimizing compiler for quantum programs.

The ISA is the key abstraction layer separating the hardware from the software, and depends heavily on the quantum computer architecture beneath. For example, for a physical quantum computer based on superconducting technology, the ISA may include controls, clock signals, and memory access.

Release News: Qiskit SDK v2.1 is here!

Technical release summary for Qiskit SDK v2.1, including updates on top new features, breaking changes, and our ongoing efforts to make Qiskit the world's most performant quantum SDK.

Qiskit v2.1 Release Notes

Today, we're excited to announce the release of Qiskit SDK v2.1! This the first major release of the Qiskit SDK v2.x series brings performance improvements and exciting new capabilities designed to enable near-term demonstrations of quantum advantage.

The dawn of quantum advantage is fast approaching, and we expect the world will need a powerful proof of practical quantum advantage by the end of 2024. Continued collaboration between the quantum and high-performance computing (HPC) communities will play an essential role in making that happen, which is why we've focused so much of our recent development efforts on extending the Qiskit SDK's C API support.

2.1

Thank you!

Any questions?



This presentation has been released under the Creative Commons CC BY-NC-ND license, i.e.

BY: credit must be given to the creator.

NC: Only noncommercial uses of the work are permitted.

ND: No derivatives or adaptations of the work are permitted.